

# Report - ADL HW2

網媒所碩一 R11944026 柯婷文

## Q1. Data Processing

### 1. Tokenizer (1%):

#### 實作上

使用這行code來呼叫hugging face上pretrained model的tokenizer。

```
tokenizer = AutoTokenizer.from_pretrained(model_args.model_name_or_path, use_fast=True)
```

由於我們的參數包括 `use_fast`，所以得到的BERT tokenizer實際上是一個基於Rust的fast tokenizer，相較於BertTokenizer更快。

透過觀察`transformers.PreTrainedTokenizer`以及`BertTokenizerFast`這兩個library的source code，可以發現實作上由 `tokenize()` 以及 `convert_tokens_to_ids()` 這兩個method實現tokenizer的主要功能：

Converts a string in a sequence of tokens (string), using the tokenizer. Split in words for word-based vocabulary or sub-words for sub-word-based vocabularies (BPE/SentencePieces/WordPieces).

步驟是

- 將現有的special token(如 `<unk>`，`<cls>` ...)跟新加入的special token(可以以參數傳入)取聯集，存在 `unique_added_tokens_encoder`
- 分別用每個special token `tok` 來切分整段文字，並將結果更新在 `text_list`；接著，用下一個 `tok` 切分更新過的 `text_list`。整個過程迭代直到所有 `tok` 都輪完為止。
- 最後再從字典裡找token對應的id，就完成了。

經過上述流程後，每個token對應的ID並會存在 `input_ids`；其他回傳值包括

- `token_type_id`，如果是兩句話中的第一句話則為0，第二句話則為1
- `attention_mask` 非padding部分為1；padding的部分為0，代表在計算attention時不用去考慮

## 觀念上

bert tokenizer主要有這兩個重要觀念：

- **WordPiece Tokenization**。做法是先把字分成subword unit，而原本是同個字的token以##作為識別，比如Sleeping就可能被切成Sleep和##ing。之後再去字典裡找token對應的ID，如此一來便可以提升token出現在vocab的機率。
- **特殊token**：針對一個由多個句子組成的文章段落，開頭會用[CLS]標示，句與句之間則用[SEP]連接。[UNK]用來表示沒出現在字典的token；padding部分用[PAD]，預訓練時使用[MASK]做MLM。這些token一樣會被轉成字典裡的ID。而 `hfl/chinese-roberta-wwm-ext` 進一步使用了 **Whole Word Masking(wwm)** 則是指用MLM任務預訓練模型時，如果word中任何一個token被mask了，那整個word的token也都會被mask。

範例如下，參考自 [中文BERT-wwm系列模型](#)

備註：在中文裡一個word指的是一個詞，character指的是一個字。

说明	样例
原始文本	使用语言模型来预测下一个词的probability。
分词文本	使用 语言 模型 来 预测 下 一个 词 的 probability。
原始Mask输入	使用 语言 [MASK] 型 来 [MASK] 测 下 一个 词 的 pro [MASK] ##lity。
全词Mask输入	使用 语言 [MASK] [MASK] 来 [MASK] [MASK] 下 一个 词 的 [MASK] [MASK] [MASK]。

## 2. Answer Span (1%):

a.

在處理QA問題時，預測結果會是一段長度不等的文字，也可說是一個區間內的文字。所以這次模型要輸出的不是token，而是token起始位置、結束位置。將資料傳給tokenizer時，設定 `return_offsets_mapping=True`，回傳值便會包括 `offsets_mapping`，紀錄的是每個word的起始、結束位置。值得注意的是，這裡會按照token原本出現在text的順序排。

text	offset_mapping
\[CLS] how many wins ...	\[(0, 0), (0, 3), (4, 8), (9, 13),...]

因為資料集提供的是answer\_start跟text兩個欄位，所以我們先用text的長度計算answer\_end的位置。如此一來便可得知答案的char-level start & end position，分別稱作 `start_char` 和 `end_char`。接著，我們要找答案的token-level start&end position。

1. 先初始化 `token_start_index` 為0、`token_end_index` 為最末端index。
2. 迭代offset\_mapping，從第一個token的offset\_mapping開始檢查，比較現在這個token的第一個char的位置(`offsets[token_start_index][0]`)是否小於等於 `start_char`，若是，就將 `token_start_index` +1，並往下一個token檢查，直到token的頭已經在 `start_char` 的右邊為止。
3. 反之，從尾部開始檢查這個token的最後一個char的位置(`offsets[token_end_index][1]`)是否大於等於 `end_char`，一路往左算回來直到token尾已經在 `end_char` 的左邊為止。
4. 最終可以得到 `token_start_index`、`token_end_index`，就是答案的token-level start&end position。

b.

首先我們可以從trainer的outputs直接得到兩個tensor分別為 `start_logits` `end_logits`，分別代表了最有可能的起始位置跟結束位置。然而因為是分開去預測的，所以有可能發生起始位置大於結束位置這種不合理的情況。於是我們分別取前 `n_best` 個logit出來，排除掉前面所說不合理的index組合後，對於所有剩下的組合，將他們機率相加，作為分數，最終取分數最高的組合，作為答案輸出。

## Q2. Modeling with BERTs and their variants

**My model - chinese-roberta-wwm-ext-large**

```
[
  {
    "_name_or_path": "hfl/chinese-roberta-wwm-ext",
    "architectures": [
      "BertForMultipleChoice"
    ],
    "attention_probs_dropout_prob": 0.1,
    "bos_token_id": 0,
    "classifier_dropout": null,
    "directionality": "bidi",
    "eos_token_id": 2,
    "hidden_act": "gelu",
    "hidden_dropout_prob": 0.1,
    "hidden_size": 768,
    "initializer_range": 0.02,
    "intermediate_size": 3072,
    "layer_norm_eps": 1e-12,
    "max_position_embeddings": 512,
    "model_type": "bert",
    "num_attention_heads": 12,
    "num_hidden_layers": 12,
    "output_past": true,
    "pad_token_id": 0,
    "pooler_fc_size": 768,
    "pooler_num_attention_heads": 12,
    "pooler_num_fc_layers": 3,
    "pooler_size_per_head": 128,
    "pooler_type": "first_token_transform",
    "position_embedding_type": "absolute",
    "torch_dtype": "float32",
    "transformers_version": "4.22.2",
    "type_vocab_size": 2,
    "use_cache": true,
    "vocab_size": 21128
  },
  {
    "_name_or_path": "hfl/chinese-roberta-wwm-ext-large",
    "architectures": [
      "BertForQuestionAnswering"
    ],
    "attention_probs_dropout_prob": 0.1,
    "bos_token_id": 0,
    "classifier_dropout": null,
    "directionality": "bidi",
    "eos_token_id": 2,
    "hidden_act": "gelu",
    "hidden_dropout_prob": 0.1,
    "hidden_size": 1024,
    "initializer_range": 0.02,
    "intermediate_size": 4096,
    "layer_norm_eps": 1e-12,
    "max_position_embeddings": 512,
    "model_type": "bert",
    "num_attention_heads": 16,
    "num_hidden_layers": 24,
    "output_past": true,
    "pad_token_id": 0,
    "pooler_fc_size": 768,
    "pooler_num_attention_heads": 12,
    "pooler_num_fc_layers": 3,
    "pooler_size_per_head": 128,
    "pooler_type": "first_token_transform",
    "position_embedding_type": "absolute",
    "torch_dtype": "float32",
    "transformers_version": "4.22.2",
    "type_vocab_size": 2,
    "use_cache": true,
    "vocab_size": 21128
  }
]
```

## Performance

train\_loss: 0.3849062052514726,  
eval\_exact\_match: 84.31372549019608,  
eval\_f1: 84.31372549019608,  
Kaggle public: 0.81029  
Kaggle private: 0.81103

## Loss function

CrossEntropyLoss()

## Optimization Algorithm

AdamW (default)

## learning rate

3e-5

## batch size

8 = (per\_device\_train\_batch\_size) 2 \* (gradient\_accumulation\_steps) 4

## Another pretrained model: macbert

```
{
  "_name_or_path": "hfl/chinese-macbert-base",
  "architectures": [
    "BertForMultipleChoice"
  ],
  "attention_probs_dropout_prob": 0.1,
  "classifier_dropout": null,
  "directionality": "bidi",
  "gradient_checkpointing": false,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.1,
  "hidden_size": 768,
  "initializer_range": 0.02,
  "intermediate_size": 3072,
  "layer_norm_eps": 1e-12,
  "max_position_embeddings": 512,
  "model_type": "bert",
  "num_attention_heads": 12,
  "num_hidden_layers": 12,
  "pad_token_id": 0,
  "pooler_fc_size": 768,
  "pooler_num_attention_heads": 12,
  "pooler_num_fc_layers": 3,
  "pooler_size_per_head": 128,
  "pooler_type": "first_token_transform",
  "position_embedding_type": "absolute",
  "torch_dtype": "float32",
  "transformers_version": "4.22.2",
  "type_vocab_size": 2,
  "use_cache": true,
  "vocab_size": 21128
}

{
  "_name_or_path": "hfl/chinese-macbert-large",
  "architectures": [
    "BertForQuestionAnswering"
  ],
  "attention_probs_dropout_prob": 0.1,
  "classifier_dropout": null,
  "directionality": "bidi",
  "gradient_checkpointing": false,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.1,
  "hidden_size": 1024,
  "initializer_range": 0.02,
  "intermediate_size": 4096,
  "layer_norm_eps": 1e-12,
  "max_position_embeddings": 512,
  "model_type": "bert",
  "num_attention_heads": 16,
  "num_hidden_layers": 24,
  "pad_token_id": 0,
  "pooler_fc_size": 768,
  "pooler_num_attention_heads": 12,
  "pooler_num_fc_layers": 3,
  "pooler_size_per_head": 128,
  "pooler_type": "first_token_transform",
  "position_embedding_type": "absolute",
  "torch_dtype": "float32",
  "transformers_version": "4.22.2",
  "type_vocab_size": 2,
  "use_cache": true,
  "vocab_size": 21128
}
```

## Performance

kaggle public: 0.8056

kaggle private: 0.80758

eval\_exact\_match: 83.5493519441675,

eval\_f1: 83.5493519441675,

train\_loss: 0.5261030852409805,

## 兩者差異

	Roberta-wwm	Macbert
Tokenization	Whole Word Masking (wwm) (更適合中文資料)	WordPiece

與Bert相比，Macbert在預訓練時引入MLM as correction的任務，讓預訓練的任務跟實際任務更為符合。macbert使用n-gram masking，並用相似詞來代替MASK，解決pretrain時有mask，而實際任務沒有的問題。

	例子
原始句子	we use a language model to predict the probability of the next word.
MLM	we use a language [M] to [M] ##di ##ct the pro [M] ##bility of the next word .
Whole word masking	we use a language [M] to [M] [M] [M] the [M] [M] [M] of the next word .
N-gram masking	we use a [M] [M] to [M] [M] [M] the [M] [M] [M] [M] [M] next word .
MLM as correction	we use a text system to ca ##lc ##ulate the po ##si ##bility of the next word .

而Roberta跟Bert的差異在於

- bert的masking是在pretrain就做了，這樣一來每個epoch的masking都一樣，而macbert使用Dynamic Masking，只在sequence送入模型中的時候才去進行動態的masking
- 移除了Bert裡面的next sentence prediction
- 使用更長的句子、更大batch size、更多資料進行訓練

### Q3. Curves


By setting `--with_tracker` option, hugging face trainer automatically exports the results to tensorboard etc. Therefore the charts below are captured from tensorboard.



eval/exact\_match  
tag: eval/exact\_match



eval/f1  
tag: eval/f1



## Q4. Pretrained vs Not Pretrained

```
{  
  "_name_or_path": "bert-base-chinese",  
  "architectures": [  
    "BertForMultipleChoice"  
  ],  
  "attention_probs_dropout_prob": 0.1,  
  "classifier_dropout": null,  
  "directionality": "bidirectional",  
  "hidden_act": "gelu",  
  "hidden_dropout_prob": 0.1,  
  "hidden_size": 768,  
  "initializer_range": 0.02,  
  "intermediate_size": 3072,  
  "layer_norm_eps": 1e-12,  
  "max_position_embeddings": 512,  
  "model_type": "bert",  
  "num_attention_heads": 12,  
  "num_hidden_layers": 12,  
  "pad_token_id": 0,  
  "pooler_fc_size": 768,  
  "pooler_num_attention_heads": 12,  
  "pooler_num_fc_layers": 3,  
  "pooler_size_per_head": 128,  
  "pooler_type": "first_token_transform",  
  "position_embedding_type": "absolute",  
  "torch_dtype": "float32",  
  "transformers_version": "4.22.2",  
  "type_vocab_size": 2,  
  "use_cache": true,  
  "vocab_size": 21128  
}
```

我將context selection的部分做調整，加入一個新的command line argument `--`



`no_pretrain`，如果有設定這個選項則在load model時不使用原本的 `from_pretrained()`，而是改用 `from_config()`，如此一來便不會load pretrained weight。值得一提的是config跟tokenizer仍然需要使用 `from_pretrained()`，因為我們還是需要模型架構跟tokenizer。

```
# old
model = AutoModelForMultipleChoice.from_pretrained(args.model_name_or_path)
# new
model = AutoModelForMultipleChoice.from_config(config)
```

the performance of this model v.s. BERT  
eval\_accuracy比較

No pretrain	Pretrain
0.5267530741110004	0.9594549684280492

## Q5. Bonus: HW1 with BERTs

```
{
  "_name_or_path": "bert-base-cased",
  "architectures": [
    "BertForSequenceClassification"
  ],
  "attention_probs_dropout_prob": 0.1,
  "classifier_dropout": null,
  "gradient_checkpointing": false,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.1,
  "hidden_size": 768,
  "initializer_range": 0.02,
  "intermediate_size": 3072,
  "layer_norm_eps": 1e-12,
  "max_position_embeddings": 512,
  "model_type": "bert",
  "num_attention_heads": 12,
  "num_hidden_layers": 12,
  "pad_token_id": 0,
  "position_embedding_type": "absolute",
  "problem_type": "single_label_classification",
  "torch_dtype": "float32",
  "transformers_version": "4.22.2",
  "type_vocab_size": 2,
  "use_cache": true,
  "vocab_size": 28996
}
```

```
{
  "_name_or_path": "bert-base-uncased",
  "architectures": [
    "BertForTokenClassification"
  ],
  "attention_probs_dropout_prob": 0.1,
  "classifier_dropout": null,
  "finetuning_task": "ner",
  "gradient_checkpointing": false,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.1,
  "hidden_size": 768,
  "initializer_range": 0.02,
  "intermediate_size": 3072,
  "layer_norm_eps": 1e-12,
  "max_position_embeddings": 512,
  "model_type": "bert",
  "num_attention_heads": 12,
  "num_hidden_layers": 12,
  "pad_token_id": 0,
  "position_embedding_type": "absolute",
  "torch_dtype": "float32",
  "transformers_version": "4.22.2",
  "type_vocab_size": 2,
  "use_cache": true,
  "vocab_size": 30522
}
```



(這張圖刪除了label2id, id2label的部分，不然貼不下)

## Performance

Intent classification

```
"eval_accuracy": 0.8756666779518127,  
"eval_loss": 1.9993304014205933,  
"train_loss": 3.157837913926171,
```

Slot tagging

```
"eval_loss": 0.09548604488372803,  
"eval_token_accuracy": 0.9679340937896072,  
"train_loss": 0.1432706798825945,
```

## Loss function

都是CrossEntropyLoss

## Optimization

都是AdamW

**lr**

5e-5 (default)

**batch\_size**

64 = s8 (per\_device\_train\_batch\_size) \* 8 (gradient\_accumulation\_steps)